MAN LUO

+1 (480) 869-7882 ◇ luoman.cs@gmail.com ◇ Website ◇ Google Scholar

PROFESSIONAL SUMMARY

Applied AI researcher with 6+ years of experience in agentic AI, multimodal retrieval and generation (RAG), synthetic data generation pipelines, foundation visionlanguage model training and evaluation, and diffusion model fine-tuning and applications. Former intern at Google, Meta, and Salesforce. Published in top-tier venues including ICML, ACL, EMNLP, AAAI, NAACL, and WACV, and author of a Springer Nature book on Multimodal Information Retrieval and Generation.

EDUCATION

Ph.D., Computer Science *August 2018 - May 2023*
Arizona State University (ASU), Tempe, Arizona, USA
Thesis: Neural Retriever-Reader for Information Retrieval and Question Answering
Chair: Dr. Chitta Baral
Committee: Dr. Yezhou Yang, Dr. Eduardo Blanco, Dr. Danqi Chen

Bachelor of Science, Computer Science *September 2014 - July 2018*
Beijing Forestry University, Beijing, China

RESEARCH AREA

Multimodal and language models with expertise in supervised fine-tuning, reinforcement learning, and evaluation; cross-modal retrieval and generation; diffusion model fine-tuning and real-world applications; scalable synthetic data generation pipelines; computer use agents for digital task automation.

INDUSTRY RESEARCH EXPERIENCE

Senior AI Research Scientist at Multimodal Cognitive Team, Intel Lab *March 2024 - Present*
Developed agentic computer use agents using reinforcement learning (DPO, GRPO with TRL), scalable synthetic data pipelines via prompt engineering, and post-training methods for foundation VLMs (2-70B) on tasks like VQA, retrieval, and instruction following. Designed diffusion-based image generation pipeline with embedded text rendering.

Research Fellow at Mayo Clinic *June 2023 - March 2024*
Built VLM and LLM pipelines for radiology report generation and longitudinal clinical note extraction using private EHR data, leveraging multimodal modeling and prompt-based fine-tuning.

Research Intern at Google Research *Aug 2022 - Dec 2022*
Designed retrieval-augmented in-context learning systems using DPR-like and BM25 retrievers to improve LLM performance on few-shot QA and multi-hop reasoning benchmarks through effective context selection and knowledge grounding.

Research Intern at Meta Reality Lab *May 2022 - Aug 2022*
Built memory-efficient hybrid retrievers and adversarial robustness evaluation frameworks for AR/VR retrieval-augmented generation (RAG) applications, combining dense-sparse retrieval and contrastive perturbation testing.

Research Intern at Salesforce.Inc *May 2021 - Aug 2021*
Evaluated generative vs. extractive QA models under domain shifts, proposing a unified benchmarking framework and deployment recommendations to guide model selection and robustness in real-world settings.

TEACHING/MENTORING

ASU Capstone Project Mentor
Detect and rewrite the toxicity in paper reviews, 5 students, *Sep 2023 - Present*

Ph.D. Mentor
Xiang Rui (Ph.D. Student at Arizona State University). *Aug 2024 - Present*
Md Messal Monem Miah (Ph.D. Student at Texas A&M University). *Oct 2023 - Present*

Master Mentor
Sanyam Lakhanpal (Master Student at ASU). *Oct 2023 - Apirl 2024*
Shrinidhi Kumbhar (Master Student at ASU). *Jan 2023 - June 2023*
Srija Macherla (SWE at Amazon). *Jan 2022 - Jun 2022*
Yankai Zeng (Ph.D student at The University of Texas at Dallas). *Aug 2020 - June 2021*

NLP Course Project Mentor
Domain Oriented Question Generation, 26 students, *Aug 2021 - Dec 2021*
Differential Diagnosis Dialogue Generation, 20 students, *Aug 2021 - Dec 2021*
Semantic Information Availability (SIA) Task, 5 students, *Jun 2020 - May 2020*
Question Answering with Varied Types of Reasoning, 5 students. *Jun 2020 - May 2020*

Teaching Assistant
CSE259 Logic in Computer Science *Dec 2020 - Dec 2021*
CSE579 Knowledge Representation and Reasoning *Aug 2019 - Dec 2019*
CSE205 Object-Oriented Programming and Data Structures *Aug 2018 - Dec 2018*

ACADEMIC SERVICE

Finance and Sponsorship Chair, WiML Workshop at NeurIPS 2025.
Editor, PLOS Digital Health.
Long-Term Super-Volunteers, WiML Workshop at NeurIPS 2024.
Organizer, Multimodal4Health at ICHI 2024
Guest Editor, PLOS Digital Medicine.
Organizer, O-DRUM at CVPR 2023.
Organizer, O-DRUM at CVPR 2022.
Reviewers, Nature Machine Intelligent, ACL, NAACL, EMNLP, EACL, AAAI, Neurips, IROS.

INVITED TALK

"Inspecting the Rise of Multimodal Through Retrieval and Content Generation Tasks" at UIUC *Oct 2024*
"Synthetic Data for Generalization and Efficiency" at ASU *Sep 2024*
"Retrieval Based In-context Learning for Large Language Models" at Google *Mar 2024*
"Advancing Multimodal Retrieval and Generation" at UMBC *Dev 2023*
"Transformer-based Multimodal Generative Model" at Mayo Clinic Radiology Showcase *Nov 2023*
"The Trend of Transformer-based Multimodal in Radiology" at RSNA *Nov 2023*
"Visual-Retriever-Reader for Knowledge-based Question Answering" at SERUM WACV *Jan 2023*
"Semantic Searching in Biomedical Domain" at exploreCSR workshop (ASU). *Mar 2021*

PUBLICATION

- Jin, H., Zhang, P., **Luo, M.** and Wang, H., 2025. Reasoning Can Hurt the Inductive Abilities of Large Language Models. Neurips 2025.
- Olson, M.L., Ratzlaff, N., Hinck, M., **Luo, M.**, Yu, S., Xue, C. and Lal, V., Semantic Specialization in MoE Appears with Scale: A Study of DeepSeek R1 Expert Specialization. EMNLP (findings) 2025.
- Su, X., **Luo, M.**, Pan, K., Chou, T., Lal, V., Howard, P. SK-VQA: Synthetic Knowledge Generation at Scale for Training Context-Augmented Multimodal LLMs. ICML 2025 (Selected as Oral, 1%) .
- Zhang, P., Jin, H., Hu, L., Li, X., Kang, L., **Luo, M.**, Song, Y. and Wang, H REVOLVE: Optimizing AI Systems by Tracking Response Evolution in Textual Optimization. ICML 2025.

- Ghaffar, U., Tariq, A., Choudry, M.M., Briggs, L.G., Channar, A., Banerjee, I., **Luo, M.**, Riaz, I.B. and Abdul-Muhsin, H.M., Domain-specific large language model for predicting prostate cancer treatment plan. Urologic Oncology: Seminars and Original Investigations 2025.
- **Luo, M.** Warren, C., Cheng, Lu., Abdul-Muhsin, H., Banerjee, I. Assessing Empathy in Large Language Models with Real-World Physician-Patient Interactions. IEEE BigData 2024.
- Lakhanpal, S., Chopra, S., Jain, V., Chadha, A., **Luo, M.** Refining Text-to-Image Generation: Towards Accurate Training-Free Glyph-Enhanced Image Generation. WACV 2025.
- **Luo, M.**, Xu, X., Liu, Y., Pasupat, P., Kazemi. In-context Learning with Retrieved Demonstrations for Language Models: A Survey. TMLR Journal 2024.
- **Luo, M.**, Trivedi, S., Kurian, A.W., Ward, K., Keegan, T.H., Rubin, D. and Banerjee, I. Automated Extraction of Patient-Centered Outcomes following Breast Cancer Treatment: An Open-Source Large Language Model-Based Toolkit. JCO Clinical Cancer Informatics 2024.
- Parmar, M., Patel, N., Varshney, N., Nakamura, M., **Luo, M.**, Mashetty, S., Mitra, A., Baral, C. Towards Systematic Evaluation of Logical Reasoning Ability of Large Language Models. ACL 2024.
- Chiang, C. C., **Luo, M.**, Dumkrieger, G., Trivedi, S., Chen, Y. C., Chao, C. J., ... & Banerjee, I. A Large Language Model-Based Generative Natural Language Processing Framework Finetuned on Clinical Notes Accurately Extracts Headache Frequency from Electronic Health Records. Headache: The Journal of Head and Face Pain 2024.
- **Luo, M.**, Xu, X., Dai, Z., Pasupat, P., Kazemi, M., Baral, C., ... Zhao, V. Y. Dr. ICL: Demonstration-Retrieved In-context Learning. NeurIPS 2023 Workshop R0-FoMo.
- **Luo, M.**, Tariq, A., Patel, B.,Banerjee, I. M3-X: Multimodal Generative Model for Screening Mammogram Reading and Explanation Medical Imaging Meets NeurIPS 2023.
- Varshney, N., **Luo, M.**, Baral, C. Exploring Training Objectives for Passage-level Differentiable Search Indexing SocalNLP 2023.
- **Luo, M.**, Tariq, A., Patel, B.,Banerjee, I. Transformer-based Multimodal Generative Model: Use-case of Screening Mammogram Reading. RSNA 2023.
- **Luo, M.** Fang, Z. Gokhale, T. Baral, C. End-to-end Knowledge Retrieval with Multi-modal Queries. ACL 2023.
- **Luo, M.**, Jain, S., Gupta, A., Einolghozati, A., Oguz, B., Chatterjee, D., Chen, X., Baral, C. and Heidari, P., 2022. A Study on the Efficiency and Generalization of Light Hybrid Retrievers. ACL 2023.
- Parmar, M., Mishra, S., Purohit, M., **Luo, M.**, Baral, C. In-BoXBART: Get Instructions into Biomedical Multi-Task Learning. NAACL 2022 Findings.
- Gokhale, T., Mishra, S., **Luo, M.**, Sachdeva, B., Baral, C. Generalized but not Robust? Comparing the Effects of Data Modification Methods on Out-of-Domain Generalization and Adversarial Robustness. ACL 2022 Findings.
- **Luo, M.**, Mitra, A., Gokhale, T., Baral, C. Improving Biomedical Information Retrieval with Neural Retrievers. AAAI 2022.
- **Luo, M.**, Zeng, Y., Banerjee, P., Baral, C. Weakly-Supervised Visual-Retriever-Reader for Knowledge-based Question Answering. EMNLP 2021.
- **Luo, M.** Sampat, S. Tallamn, R. Zeng, Y. Vancha, M. Sajja, A. Baral, C. Just because you are right, doesnt mean I am wrong: Overcoming a bottleneck in development and evaluation of Open-Ended VQA tasks. EACL 2021.
- Lee, J. and **Luo, M.**, 2019. Strong equivalence for LPMLN programs. ICLP 2019.

- Varshney, N., **Luo, M.**, Baral, C. Can Open-Domain QA Reader Utilize External Knowledge Efficiently like Humans? AAAI 2023 Workshop on Knowledge Augmented Methods for NLP
- **Luo, M.**, Parmar, M., Mahendran, J. S., Jain, S., Rawal, S., Baral, C. SCONER: Scoring Negative Candidates Before Training Neural Re-Ranker For Question Answering ICML 2022 Workshop on Knowledge Retrieval and Language Models.
- **Luo, M.**, Saxena, S., Mishra, S., Parmar, M., Baral, C. BioTABQA: Instruction Learning for Biomedical Table Question Answering CEUR Workshop 2022.
- **Luo, M.** Neural Retriever and Go Beyond: A Thesis Proposal. NAACL 2022 Student Research Workshop.
- **Luo, M.**, Chen, S., Baral, C A Simple Approach to Jointly Rank Passages and Select Relevant Sentences in the OBQA Context NAACL 2022 Student Research Workshop.
- **Luo, M.**, Hashimoto, K., Yavuz, S., Liu, Z., Baral, C., Zhou, Y. Choose Your QA Model Wisely: A Systematic Study of Generative and Extractive Readers for Question Answering ACL 2022 Spa-NLP workshop.

PRE-PRINT

- **M Luo**, B Peterson, R Gan, H Ramalingame, N Gangrade, A Dimarogona I Banerjee, P Howard Benchmark on Peer Review Toxic Detection: A Challenging Task with a New Dataset arXiv preprint 2025.
- E Aflalo, GBM Stan, T Le, **M Luo**, S Rosenman, S Paul, SY Tseng, V Lal FiVL: A Framework for Improved Vision-Language Alignment arXiv preprint 2025.
- N Ratzlaff, **M Luo**, X Su, V Lal, P Howard Training-Free Mitigation of Language Reasoning Degradation After Multimodal Instruction Tuning arXiv preprint 2025.
- G Ben-Melech Stan, E Aflalo, **M Luo**, S Rosenman, T Le, S Paul, SY Tseng, V Lal FastRM: An efficient and automatic explainability framework for multimodal generative models arXiv preprint 2025.
- PK Choubey, X Su, **M Luo**, X Peng, C Xiong, T Le, S Rosenman, V Lal, P Mui, R Ho, P Howard, CS Wu Distill-SynthKG: Distilling Knowledge Graph Synthesis Workflow for Improved Coverage and Efficiency arXiv preprint 2024.
- **Luo, M.**, Kumbhar, S., Parmar, M., Varshney, N., Banerjee, P., Aditya, S., Baral, C. Towards LogiGLUE: A Brief Survey and A Benchmark for Analyzing Logical Reasoning Capabilities of Language Models. arXiv preprint 2023.
- Macherla, S., **Luo, M.**, Parmar, M., Baral, C. MDDial: A Multi-turn Differential Diagnosis Dialogue Dataset with Reliability Evaluation. arXiv preprint 2023.
- **Luo, M.**. Neural Retriever-Reader for Information Retrieval and Question Answering (Doctoral dissertation, Arizona State University, 2023).
- Varshney, N., Parmar, M., Patel, N., Handa, D., Sarkar, S., **Luo, M.**, Baral, C.. Can NLP Models Correctly Reason Over Contexts that Break the Common Assumptions?. arXiv preprint 2023.
- Liu, Z., Chen, Y., Li, J., **Luo, M.**, Yu, P. S., Xiong, C. Improving contrastive learning with model augmentation. arXiv preprint 2022.
- Banerjee, P., Baral, C., **Luo, M.**, Mitra, A., Pal, K., Son, T. C., Varshney, N. Can Transformers Reason About Effects of Actions? arXiv preprint, 2020.

BOOK MANUSCRIPT

Luo, M. Gokhale, T., Varshney, N., Yang, Y., Baral, C. Advances in Multi-Modal Information Retrieval and Generation. *Springer Nature*

AWARD

Finalist of 2021 Knowledge Mobilization Awards. Website *April 2021*

2019 ICLP conference Doctoral Consortium Travel Award. Website *September 2019*

Honorable Mention in Interdisciplinary Contest in Modeling(ICM) *April 2017*